

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

COMP

# Review (1/2)

- Caches are NOT mandatory:
  - Processor performs arithmetic
  - Memory stores data
  - Caches simply make t
- Each level of memory hierarchy is a set of next higher level
- Caches speed up due to **temporal locality**: store data used recently
- Block size > 1 word speeds up due to **spatial locality**: store words adjacent to the ones used recently

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

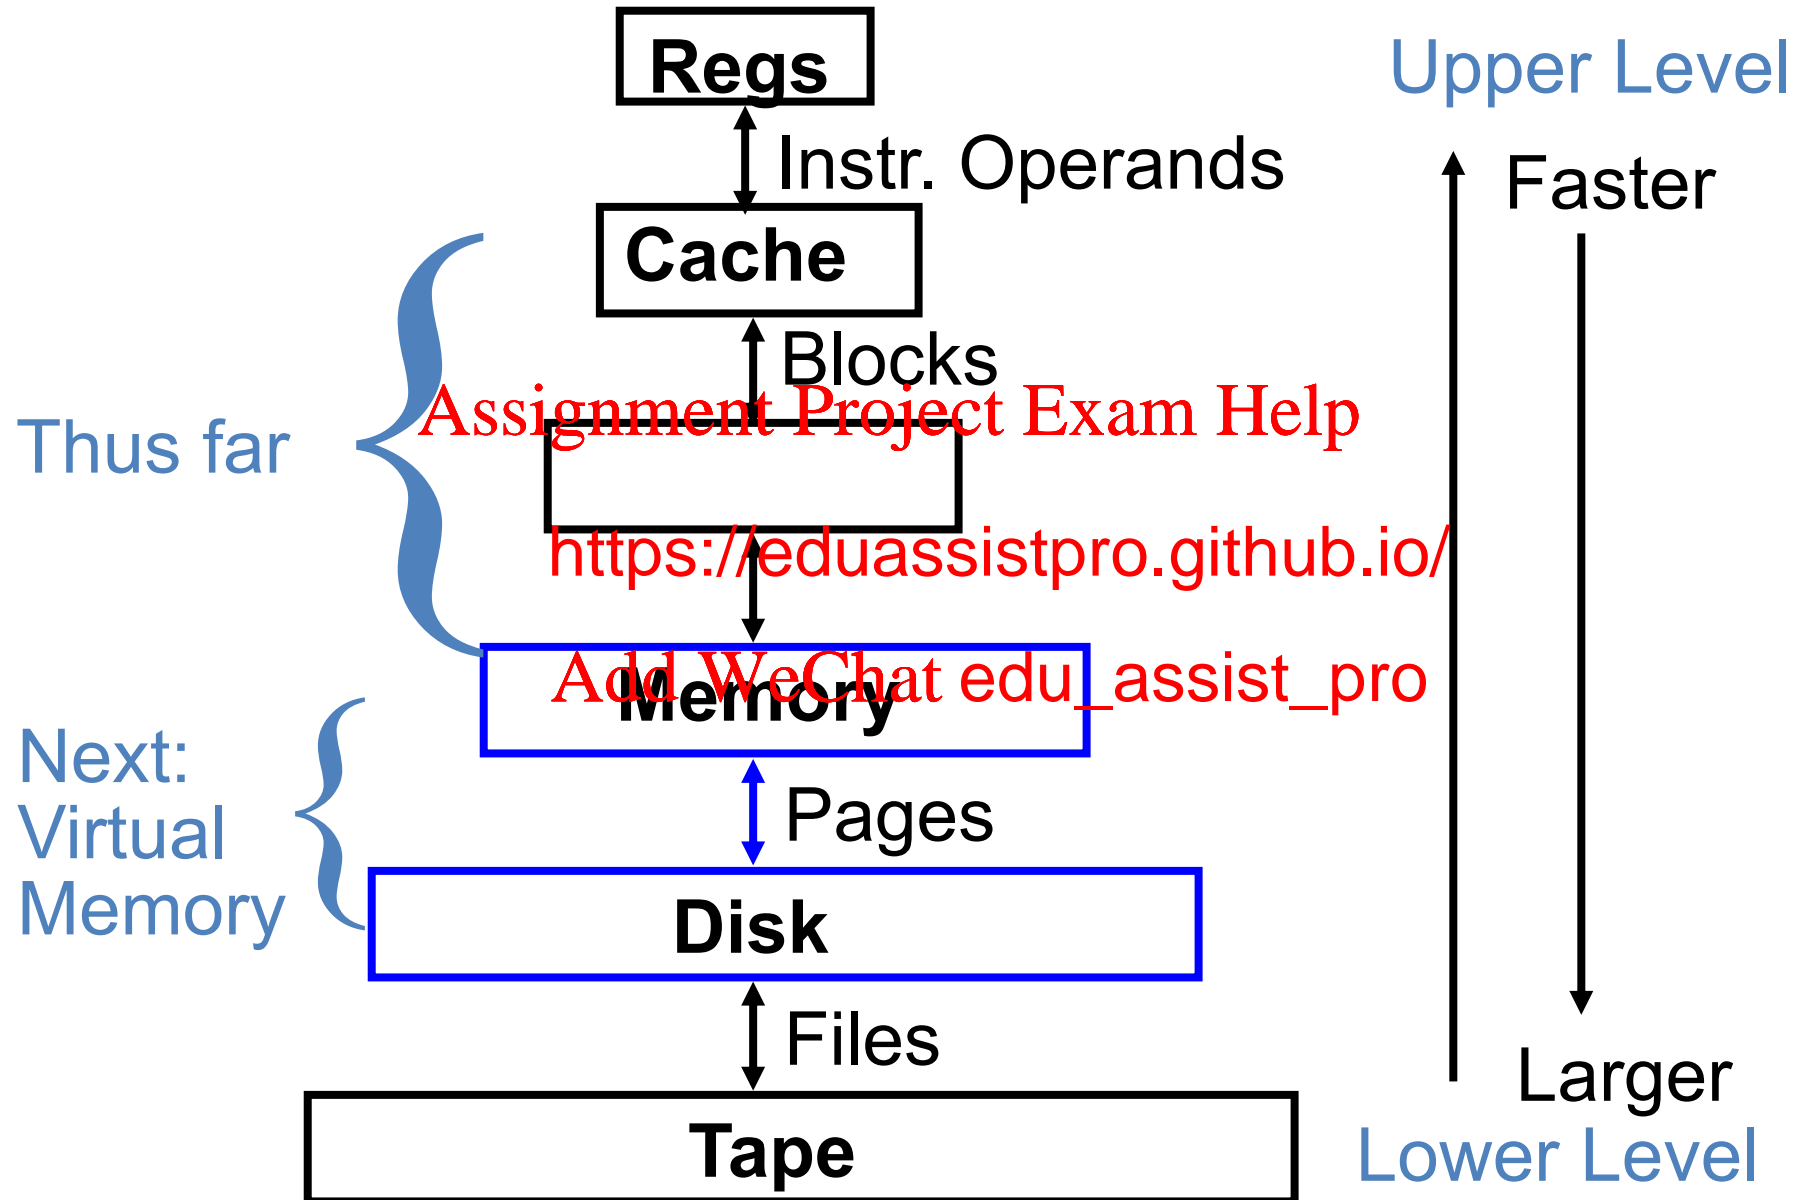
# Review (2/2)

- Cache design choices:
  - size of cache: speed v. capacity
  - direct-mapped v. associative
  - for N-way set assoc: <https://eduassistpro.github.io/>
  - block replacement policy
  - 2nd level cache?
  - Write through v. write back?
- Use performance model to pick between choices, depending on programs, technology, budget, ...

Assignment Project Exam Help

Add WeChat edu\_assist\_pro

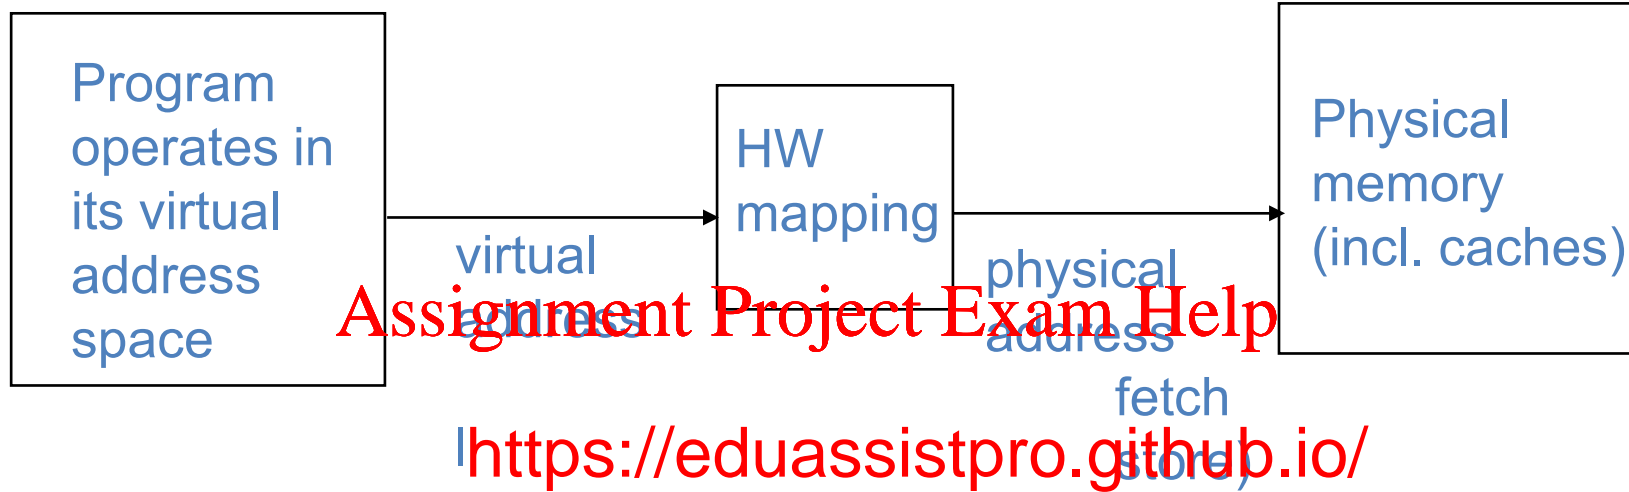
# Another View of the Memory Hierarchy



# Virtual Memory

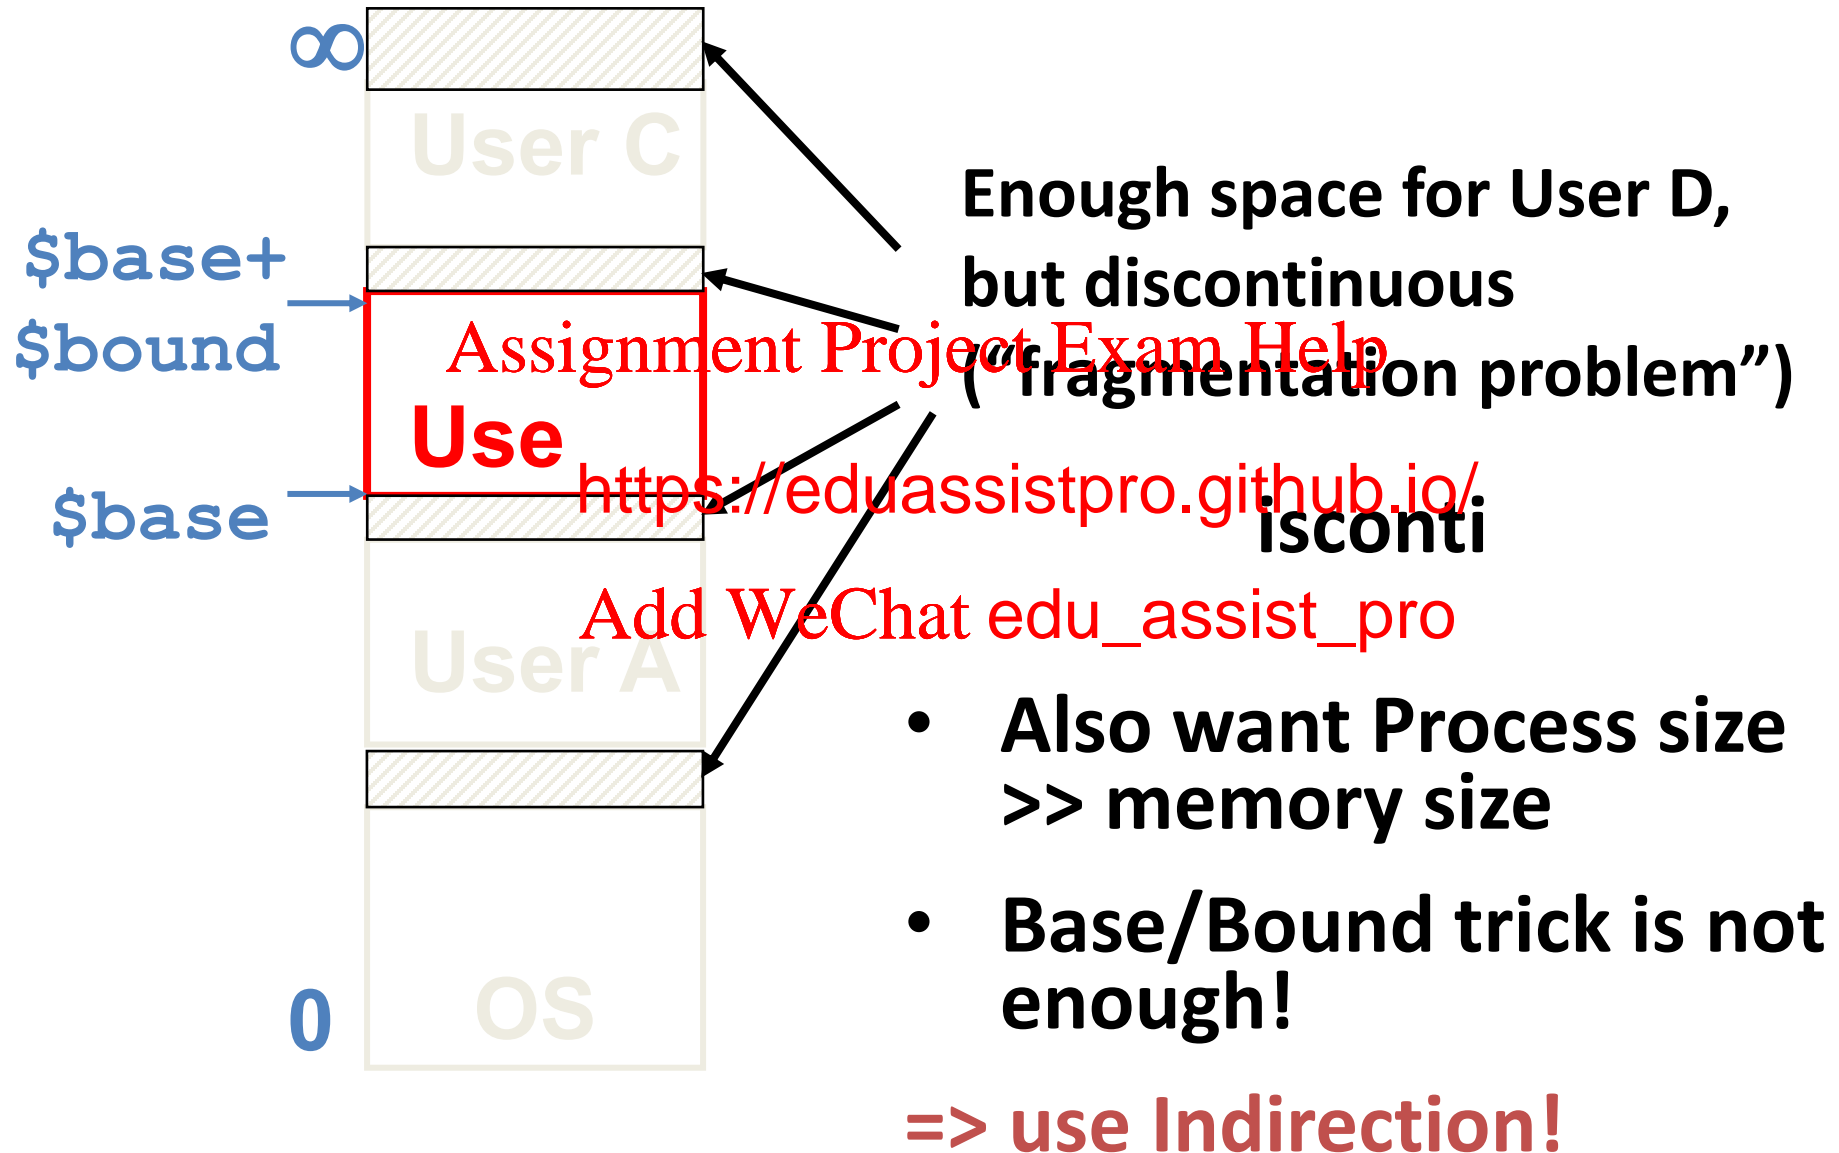
- If Principle of Locality allows caches to offer (usually) speed of cache memory with size of DRAM memory, then recursively why not use at next level to give speed of DRAM memory, size of Disk <https://eduassistpro.github.io/>
- Called “Virtual Memory”  
Add WeChat edu\_assist\_pro
  - Also allows OS to share memory, protect programs from each other
  - Today, more important for protection vs. just another level of memory hierarchy
  - Historically, it predates caches

# Virtual to Physical Addr. Translation



- Each program operates in its own virtual address space; as if it were the only program running
- Each “process” is protected from the other
- OS can decide where each goes in memory
- Hardware (HW) provides virtual -> physical mapping

# Simple Example: Base and Bound Reg




# Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

## **GDT: Global descriptor table**

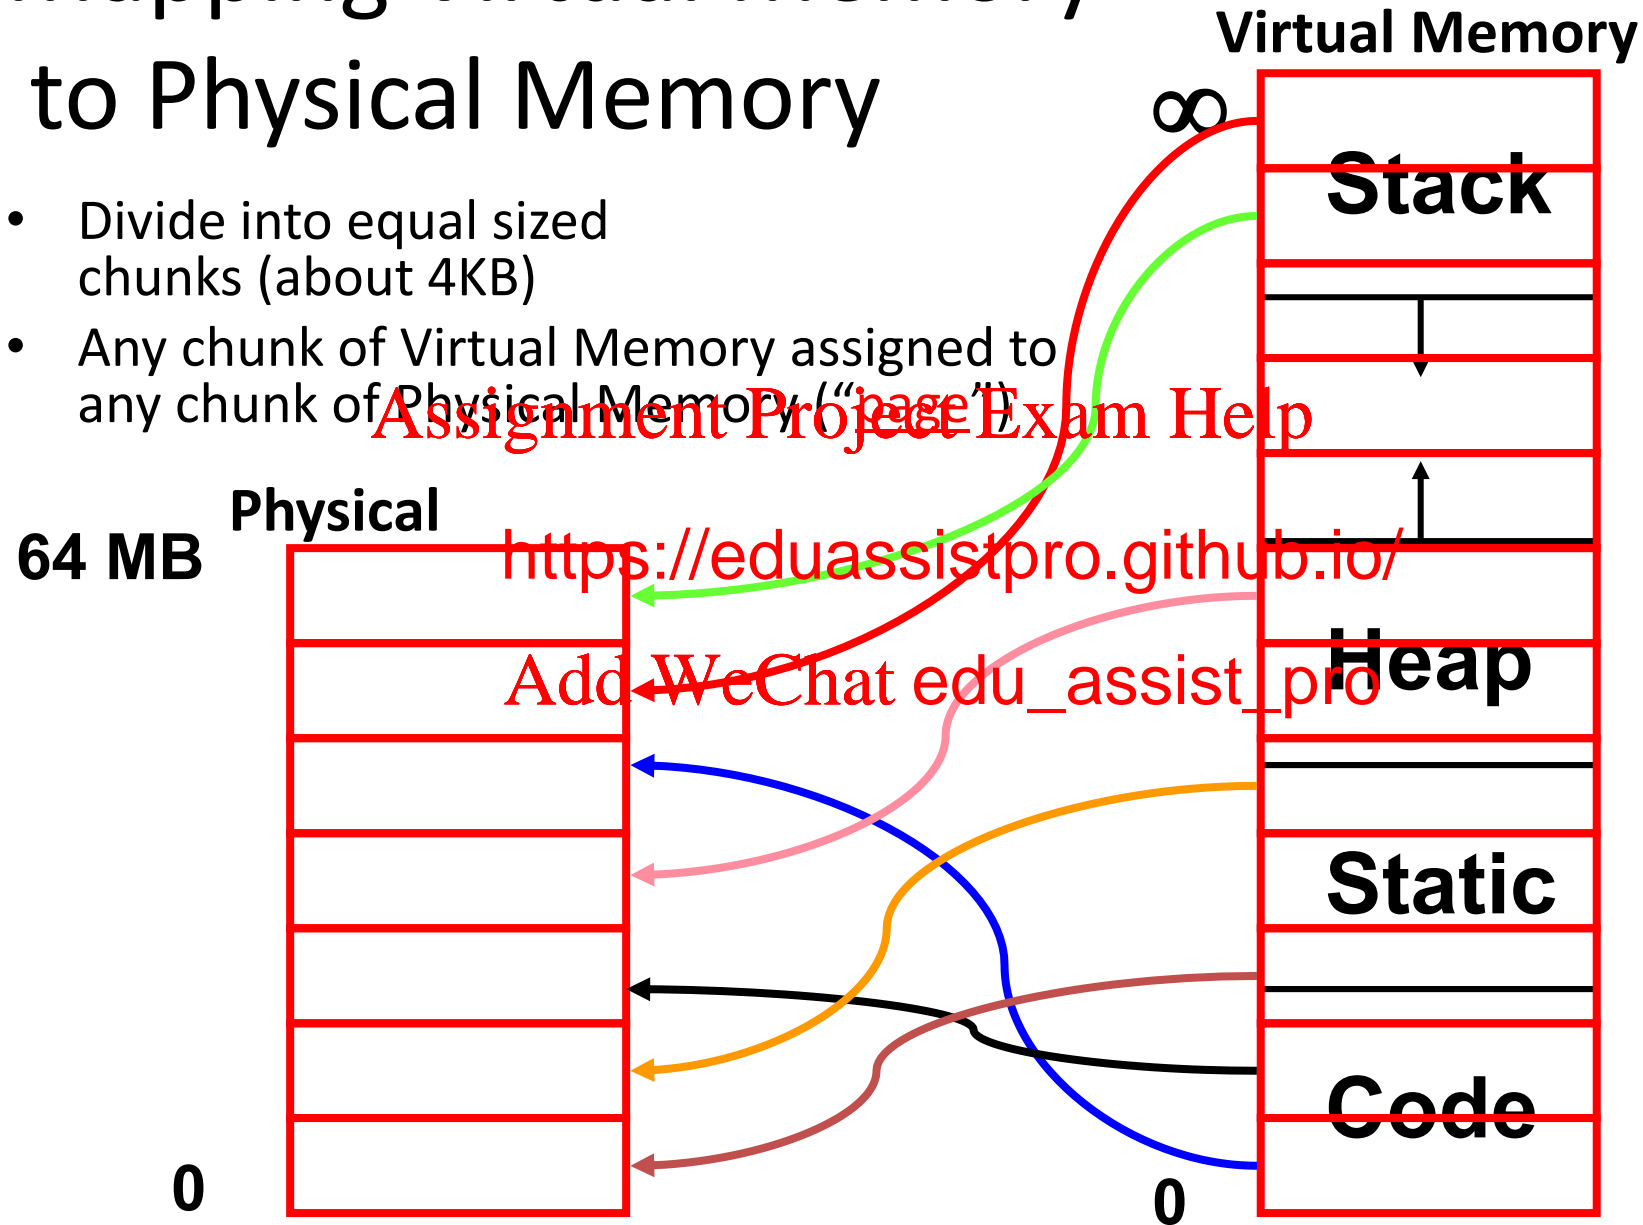
On 286 machines, it allowed for base and bound, while indirection and virtual memory was not introduced to x86 chips until the 386 (in 1986)





# Mapping Virtual Memory to Physical Memory

- Divide into equal sized chunks (about 4KB)
- Any chunk of Virtual Memory assigned to any chunk of Physical Memory ("page")



# Virtual Memory Mapping Function

- Cannot have simple function to predict arbitrary mapping
- Use table lookup of mappings

**Virtual address:**

Page Number	Offset
-------------	--------

- Use table lookup <https://eduassistpro.github.io/>
  - Page number is index
- Virtual Memory Mapping Function
  - Physical Offset = Virtual Offset
  - Physical Page Number  
= PageTable[Virtual Page Number]
  - (P.P.N. also called "[Page Frame](#)")

# Page Table

- A page table is an operating system structure which contains the mapping of virtual addresses to physical locations
  - There are several different ways, all up to the operating system, to keep this data around
- Each process running in the operating system has its own page table
  - “State” of process is PC, all registers, plus page table
  - OS changes page tables by changing contents of Page Table Base Register

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Address Mapping: **Page Table**

Virtual Address:

**page no.**   **offset**

Page Table  
Base Reg

index  
into  
page  
table

**Page Table**  
Assignment Project Exam Help

Valid	Access Rights	Address
V	A.R.	P. P. A.
V	A.R.	P. P. A.
...		

<https://eduassistpro.github.io/>  
Add WeChat edu\_assist\_pro

II

(concatenation)

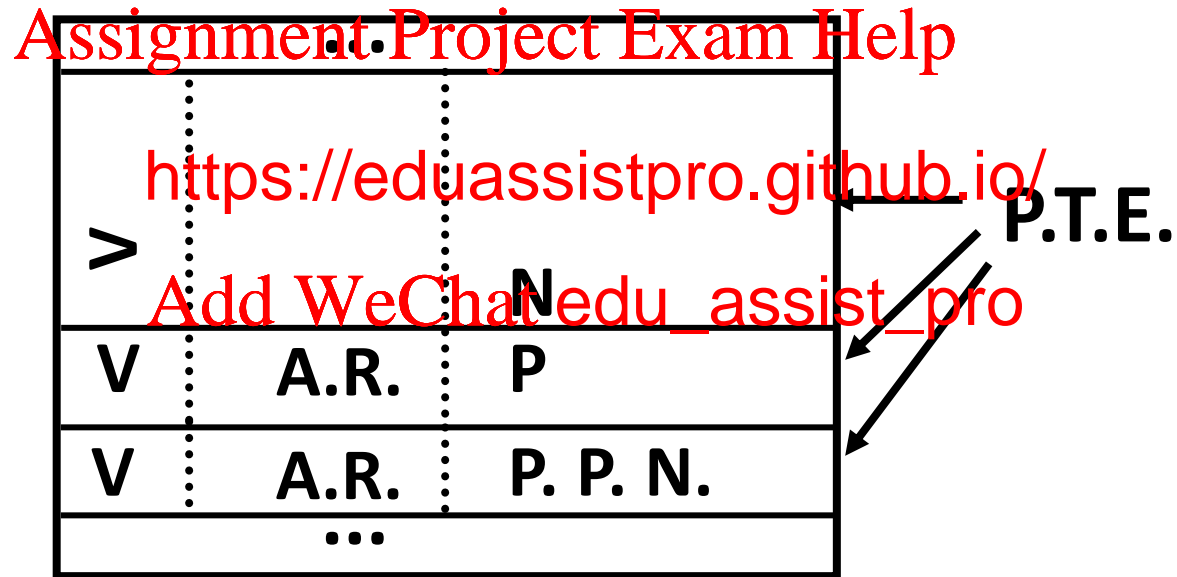
Physical  
Memory  
Address

Page Table located in physical memory

# Page Table Continued

- Page Table Entry (PTE) Contains either Physical Page Number or indication not in Main Memory (Valid = 0)
  - OS maps to disk if Not Valid (V = 0)

Page Table



- If valid, check permission to use page: **Access Rights** (A.R.) may be Read Only, Read/Write, Executable

# Notes on Page Table

- Solves Fragmentation problem: all chunks same size, so all holes can be used
- OS must reserve “Swap Space” on disk for each process
- To grow a program
  - If unused pages
  - If not, OS swaps some old pages
  - (Least Recently Used to pick pages to swap)
- Each process has own Page Table
- Will add details, but Page Table is essence of Virtual Memory

# Analogy

- Book title like **virtual address** “See MIPS Run”
- Call number like **physical address** QA76.9 A73 S88 2007
- Card catalogue like **page number** from book title to call number  
<https://eduassistpro.github.io/>
- On card for book, in local library or **remote branch** like **valid bit** indicating in main memory or on disk
- On card, loan restrictions are like **access rights**, for instance, reserved for 2-hour in library use, or 2-week checkout

# Comparing the 2 levels of hierarchy

## Cache Version

Block (or Line)

Miss

Block Size: 32-64B

Placement:

Direct Mapped,  
N-way Set Associative

Replacement:  
LRU or Random or...

Write Thru or Back

## Virtual Memory Version

Page

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

Least Recently Used  
(LRU)

Write Back

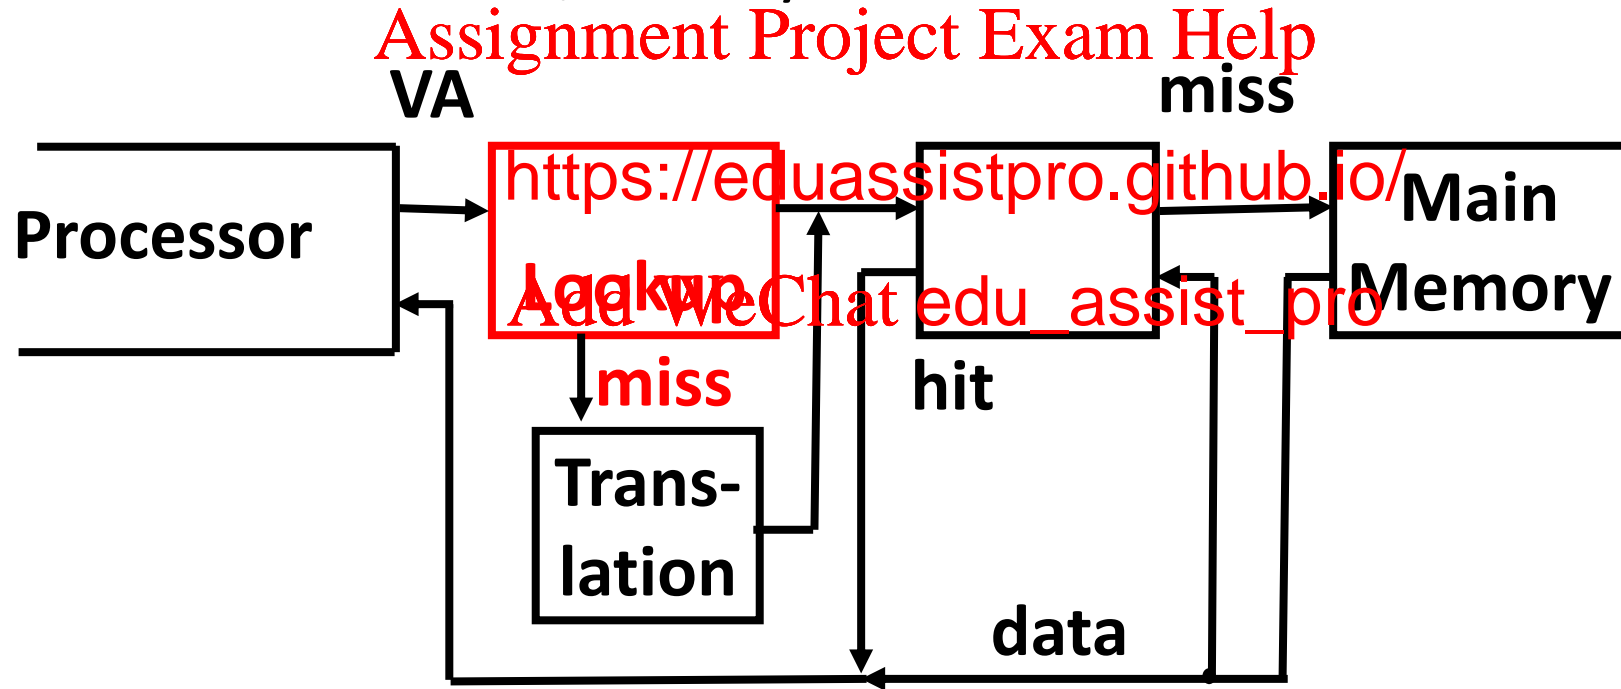


# Virtual Memory Problem #1

- Map every address  $\Rightarrow$  1 indirection via Page Table in memory  
per virtual address  $\Rightarrow$  1 virtual memory access  
– This implies 2 physical memory accesses **SLOW!**
- Observation: since I <https://eduassistpro.github.io/> data, there must be locality in **virtual address translation** these pages
- Since small is fast, why not use a small cache of virtual to physical address translations to make translation fast?
- For historical reasons, this cache is called a **Translation Lookaside Buffer**, or **TLB**

# Translation Look-Aside Buffers

- TLBs usually small, typically 128 - 256 entries
- Like any other cache, the TLB can be direct mapped, set associative, or fully associative



# Typical TLB Format

Virtual Address	Physical Address	Dirty	Ref	Valid	Access Rights

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- TLB just a cache
- TLB access time comparable to main memory access time (much less than main memory access time)
- Dirty: since we use “write back” policy, need to know whether or not to write page to disk when replaced
- Ref: Used to help calculate LRU on replacement
  - Can be cleared periodically by OS, then checked later to see if page was **referenced**

# What if page not in TLB?

- Option 1: Hardware checks page table and loads new Page Table Entry into TLB
- Option 2: Hardware <https://eduassistpro.github.io/> S to decide what to do
  - This is a simple and fast
- MIPS follows Option 2. Hardware [Add WeChat edu\\_assist\\_pro](#) nothing about page table
  - That is, there is no “page table base register” and instead there is only the TLB

# TLB Miss (simple strategy)

- If the address is not in the TLB, MIPS traps to the operating system
  - When in the operating system, we don't do translation (turn off virtual memory)
- The operating system knows which virtual address desired was requested
  - So we look the entry up in the page table
  - Then we add the entry to the TLB
  - Then we resume the program again at the instruction that failed (it will not have a TLB miss next time)

# What if the data is on disk?

- We load the page off the disk into a free block of memory, using a **DMA transfer**
  - Meantime we switch to another process waiting to be run
- When the DMA is complete, the process is interrupted, and can then update the process's page table
  - So when we switch back to the task, the desired data will be in memory

Assignment Project Exam Help

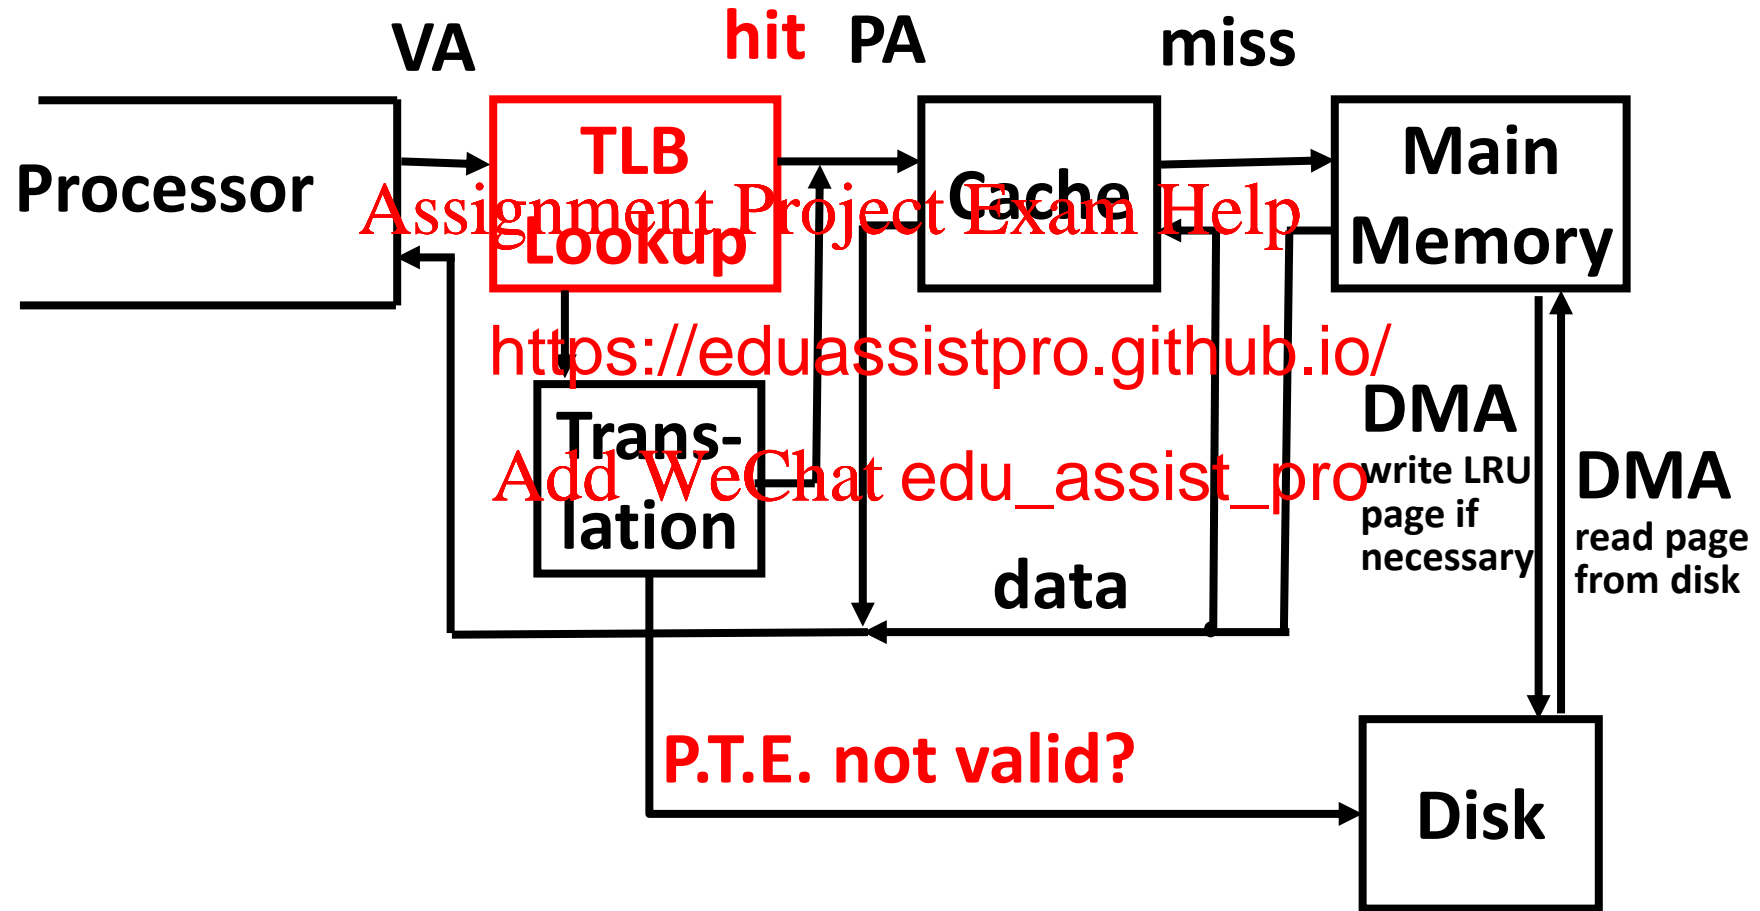
<https://eduassistpro.github.io/>

Add WeChat: edu\_assist\_pro

# What if we don't have enough memory?

- Chose some physical page belonging to a program,
  - We chose the physical page to evict based on replacement policy (e.g., LRU)
  - If chosen page is not on disk, write it to the disk
  - If chosen page is on disk but not up-to-date, then we can simply overwrite it in memory
- The program previously using the chosen page must have its page table updated to reflect the fact that its memory moved somewhere else.
- Finally, the OS can update our program's page table to use this physical page

# Data on Disk?





# Virtual Memory Problem #2



- Not enough physical memory! Suppose 4KB pages and...
  - Have only 64 MB ( $2^{26}$  B) of physical memory
  - N processes, each given 1 GB ( $2^{30}$  B) of virtual memory
  - How many virtual pages?  $N=1$  process?
  - For what N will we have 1 physical page?
- Spatial Locality to the rescue
  - Each page is 4 KB, lots of nearby references
- Even for huge programs, typically only accessing a few pages at any given time
  - The “Working Set” of recently used pages

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Virtual Memory Problem #3

- Page Table too big!

- 4 GB Virtual Memory  $\div$  4 KB page

- $\Rightarrow$  approximately 1 million Page Table Entries,

- each taking up 1

- $\Rightarrow$  4 MB just for Pa

- 25 processes will need 100 MB Tables!

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

- Variety of solutions to tradeoff memory size of *mapping function* for *slower TLB misses*

- Make TLB large enough, highly associative so rarely miss on address translation

- *Alternative mapping functions are not in the scope of this course*

# Things to Remember 1/2

- Apply Principle of Locality Recursively
- Reduce Miss Penalty? add a (L2) cache
- Manage memory to
  - Originally included protection as b protection is critical
  - Use Page Table of mappings vs. tag/data in cache
- Virtual memory to Physical Memory Translation too slow?
  - Add a cache of Virtual to Physical Address Translations, called a TLB

# Things to Remember 2/2

- Virtual Memory allows protected sharing of memory between processes with less swapping to disk, less fragmentation than always-swap, or base/bound
- Spatial and Temporal <https://eduassistpro.github.io/> of Pages is all that must be in memory for process to well
- TLB to reduce performance cost of VM
- Need a more compact representation to reduce memory size cost of simple 1-level page table (*especially when using a 64-bit address space*)

# Review and More Information

- Textbook 5<sup>th</sup> edition 5.7, Virtual Memory

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro